

Data Structures: Vectors

Enrico Toffalini

PSICOSTAT

What Are Data Structures

Data structures, like **vectors**, **matrices**, **dataframes**, **lists**, are fundamental tools that allow you to **organize** and **store complex information**, so that they can be easily **processes by functions** (e.g., `lm()` function may fit a linear model on variables stored in a dataframe)

Most operations you will perform in R (e.g., *processing data, fitting models, plotting outputs*) are performed on these data structures

Vectors

Simple one-dimensional structures that store data of different types

<i>Position:</i>	[1]	[2]	[3]	[4]	...	[n]
<i>Value:</i>	X_1	X_2	X_3	X_4	...	X_n

Here is an actual **example** (of a *numerical* vector):

<i>Position:</i>	[1]	[2]	[3]	[4]	...	[n]
Hours	10	15	20	10	...	5

Vectors as 1-D Arrays

Vectors are just special cases of arrays

Scalar
0-Dimensional array

0.5

Vectors
1-Dimensional array

6	9.7	2	3.1	6.73	4
---	-----	---	-----	------	---

6 elements

Matrix
2-Dimensional array

9	6.5	3.4	0	0.2	12
10	73.2	8	2	0.4	11
34	52	19	1.6	4	15

3 rows

6 columns

3-Dimensional Array

12	4	2.1	4.5	8
0.3	19	6	0.9	11
10	7.4	0.1	91	4.3
84	0.5	15	81	5

4 rows

5 columns

3 slices

+

Create Vectors with `c()`

Vectors can easily be created using the `c()` base function, with a sequence of elements separated by commas “,”

Vectors can be of different types. The following example shows a *character* vector (note the quotes " " around objects):

```
Teachers = c("Pastore", "Kiesner", "Granziol", "Toffalini",
            "Calignano", "Epifania", "Bastianelli")
```

or numeric:

```
Hours = c(10, 15, 20, 10, 15, 5, 15, 5)
```

Vectors Must be Homogeneous

Vectors must contain elements of the **same type**. If you mix types, R will automatically **coerce** the elements to a single type, which may lead to undesired results.

Therefore, **avoid mixing data types!** Example:

```
Hours = c(10, 15, 20, 10, 15, "tbd", 15, 5)  
Hours  
[1] "10"  "15"  "20"  "10"  "15"  "tbd" "15"  "5"
```

everything was coerced to become a character!

If needed, use **NA** (Not Available):

```
Hours = c(10, 15, 20, 10, 15, NA, 15, 5)  
Hours # remains a numerical vector, NA does not affect type  
[1] 10 15 20 10 15 NA 15  5
```

Vectors Must be Homogeneous

You may coerce a vector to be a particular type if needed

```
Hours = c(10, 15, 20, 10, 15, "tbd", 15, 5)
Hours
[1] "10"  "15"  "20"  "10"  "15"  "tbd" "15"  "5"
as.numeric(Hours)
```

Warning: NAs introduced by coercion

```
[1] 10 15 20 10 15 NA 15  5
```

But be careful! Elements that cannot be coerced to the target type, will be replaced with NA

```
Hours = c("10", "15,", "20", " 10", "15 ", "tbd", "15.", "5_")
as.numeric(Hours)
```

Warning: NAs introduced by coercion

```
[1] 10 NA 20 10 15 NA 15 NA
```

Indexing Vectors

Select/extract elements with INDEXING using square brackets []:

```
Hours = c(10, 15, 20, 10, 15, 5, 15, 5)
```

```
Hours[4] # a single element
```

```
[1] 10
```

```
Hours[5:7] # a range of elements
```

```
[1] 15 5 15
```

```
Hours[c(1,3,6)] # specific elements
```

```
[1] 10 20 5
```

Know the length of a vector using the length() function, and use it:

```
length(Hours)
```

```
[1] 8
```

```
Hours[length(Hours)] # use it to extract the last element
```

```
[1] 5
```

Indexing Vectors

Negative indexing

You can use the *minus sign -* to select **all elements except some** from a vector.
(This method is also applicable to dataframes)

```
Hours = c(10, 15, 20, 10, 15, 5, 15, 5)
```

```
Hours[-4] # ALL BUT a single element
```

```
[1] 10 15 20 15 5 15 5
```

```
Hours[-c(5:7)] # ALL BUT a range of elements
```

```
[1] 10 15 20 10 5
```

```
Hours[-c(1,3,6)] # ALL BUT specific elements
```

```
[1] 15 10 15 15 5
```

```
Hours[-length(Hours)] # ALL BUT the last element
```

```
[1] 10 15 20 10 15 5 15
```

Logical Indexing

Often, you'll need to extract values from a vector based on specific *logical* conditions. Here's an example:

```
Hours = c(10, 15, 20, 10, 15, 5, 15, 5)
Hours[Hours >= 15] # extract only values greater than or equal to 15
[1] 15 20 15 15
```

This is called *logical indexing* because you are selecting elements based on a logical vector (i.e., a sequence of `TRUE`, `FALSE`):

```
Hours >= 15 # the logical vector actually inside the square brackets
[1] FALSE TRUE TRUE FALSE TRUE FALSE TRUE FALSE
```

Also, you can use a vector to extract values from another vector:

```
Teachers[Hours >= 15]
[1] "Kiesner"      "Granzio1"     "Calignano"    "Bastianelli"
```

Indexing and Assignment

With indexing, you can not only select, but also **assign** or **modify** elements in a vector:

```
Hours = c(10, 15, 20, 10, 15, 5, 15, 5)
```

```
Hours[1] = 0 # assign a new value
```

```
Hours[3] = Hours[3]+50 # modify an existing element
```

```
Hours
```

```
[1] 0 15 70 10 15 5 15 5
```

You can even assign values **outside the current range** of the vector. But what happens?

```
Hours[20] = 5
```

```
Hours
```

```
[1] 0 15 70 10 15 5 15 5 NA 5
```

Operating on Vectors

you can simultaneously apply an operation to a whole vector, like

```
Hours = c(10, 15, 20, 10, 15, 5, 15, 5)
```

```
Hours / 5
```

```
[1] 2 3 4 2 3 1 3 1
```

Of course, this is useful when you want to save the result as a new vector:

```
ECTS = Hours / 5
```

Similarly, you can apply functions to all elements of a vector:

```
sqrt(Hours) # computes square root of each element
```

```
[1] 3.162278 3.872983 4.472136 3.162278 3.872983 2.236068 3.872983 2.236068
```

```
log(Hours) # computes the natural Logarithm of each element
```

```
[1] 2.302585 2.708050 2.995732 2.302585 2.708050 1.609438 2.708050 1.609438
```

Summary Statistics on Vectors

A whole vector may serve to compute summary statistics, for example using functions such as `mean()`, `sd()`, `median()`, `quantile()`, `max()`, `min()`:

```
mean(Hours) # returns the average value (mean) of the vector
```

```
[1] 11.875
```

```
sd(Hours) # returns the Standard Deviation of the vector
```

```
[1] 5.303301
```

```
median(Hours) # returns the median value of the vector
```

```
[1] 12.5
```

Summary Statistics on Vectors

A whole vector may serve to compute summary statistics, for example using functions such as `mean()`, `sd()`, `median()`, `quantile()`, `max()`, `min()`:

```
quantile(Hours, probs=c(.25, .50, .75)) # returns desired quantiles
```

```
25%   50%   75%
8.75 12.50 15.00
```

```
max(Hours) # returns largest value
```

```
[1] 20
```

```
min(Hours) # returns smallest value
```

```
[1] 5
```

Summary Statistics - Managing Missing (**NA**) Values

All of the previous summary statistics will **fail** if there is even a single **NA** value:

```
Hours = c(10, 15, 20, 10, 15, NA, 15, 5)
```

```
mean(Hours) # a single NA value implies that the average is impossible to determine
```

```
[1] NA
```

```
quantile(Hours, probs=c(.25, .75)) # quantile() will even return an Error
```

Error in quantile.default(Hours, probs = c(0.25, 0.75)): missing values and
NaN's not allowed if 'na.rm' is FALSE

You can easily manage missing values by adding the **na.rm=TRUE** argument:

```
mean(Hours, na.rm=TRUE) # NA values are ignored
```

```
[1] 12.85714
```

```
quantile(Hours, probs=c(.25, .75), na.rm=TRUE) # NA values are ignored
```

```
25% 75%
```

```
10 15
```

Replacing NA With the Average Value

Replacing a missing value with the average across valid values is risky, as it may alter many other summary statistics, but it is a good example for understanding different concepts seen so far:

```
Hours = c(10, 15, 20, 10, 15, NA, 15, 5)
```

```
# compute the average value ignoring NAs, and put it wherever  
# there is a NA value in the vector
```

```
Hours[is.na(Hours)] = mean(Hours, na.rm=TRUE)
```

```
# now let's inspect the updated content of the vector
```

```
Hours
```

```
[1] 10.00000 15.00000 20.00000 10.00000 15.00000 12.85714 15.00000 5.00000
```

```
# by the way... na.rm=TRUE is no longer needed now, as NA is no longer there  
mean(Hours)
```

```
[1] 12.85714
```

Frequency Counts

Another useful summary statistic is the **frequency count**, which shows how often each unique value appears in a vector. You can use the **table()** function to calculate frequencies easily:

```
type = c("METHODOLOGY", "METHODOLOGY", "PROGRAMMING", "SOFT SKILLS", "SOFT SKILLS",
        "METHODOLOGY", "SOFT SKILLS", "METHODOLOGY", "PROGRAMMING")
table(type)
```

```
type
METHODOLOGY PROGRAMMING SOFT SKILLS
        4             2            3
```

Be careful: R is case sensitive!

```
type = c("METHODOLOGY", "methodology", "PROGRAMMING", "SOFT SKILLS", "SOFT SKILLS",
        "METHODOLOGY", "SOFT SKILLS", "METHODOLOGY", "Programming")
table(type)
```

```
type
methodology METHODOLOGY Programming PROGRAMMING SOFT SKILLS
        1             3             1             1            3
```